

MUQ: The MIT Uncertainty Quantification Library

Matthew Parno^{*1, 2}, Andrew Davis³, and Linus Seelinger⁴

¹ US Army Corps of Engineers, Cold Regions Research and Engineering Laboratory, Hanover, NH
² USA ² Dartmouth College, Hanover, NH USA ³ Courant Institute of Mathematical Sciences, New
⁴ York University, New York, NY USA ⁴ Institute for Scientific Computing, Heidelberg University,
Heidelberg, Germany

DOI: [10.21105/joss.03076](https://doi.org/10.21105/joss.03076)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Pierre de Buyl](#) ↗

Reviewers:

- [@martinmodrak](#)
- [@georgiastuart](#)

Submitted: 26 February 2021

Published: 03 March 2021

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Scientists and engineers frequently rely on mathematical and numerical models to interpret observational data, forecast system behavior, and make decisions. However, unknown and neglected physics, limited and noisy data, and numerical error result in uncertain model predictions. The MIT Uncertainty Quantification library (MUQ) is a modular software framework for defining and solving uncertainty quantification problems involving complex models. MUQ is written in c++ but uses pybind11 ([Jakob et al., 2017](#)) to provide a nearly comprehensive python interface. Users can access nearly all of MUQ's capabilities from either language.

MUQ provides users many commonly used UQ tools and its modular design allows developers to easily modify, extend, and advance existing algorithms. For example, MUQ allows exact sampling of non-Gaussian distributions (e.g., Markov chain Monte Carlo and importance sampling), approximating computationally intensive forward models (e.g., polynomial chaos expansions and Gaussian process regression), working with integral covariance operators (e.g., Gaussian processes and Karhunen-Loève decompositions), and characterizing predictive uncertainties. The software is designed to support algorithm developers who want to easily construct new algorithms by exploiting a wide variety of existing algorithmic building blocks. Many UQ algorithms are model agnostic: Different physics-based or statistical models can be substituted into the algorithm based on the application. Therefore, MUQ enables users to quickly implement new models and exploit state-of-the-art UQ algorithms.

A suite of documented examples, including Gaussian process regression of Mauna Loa CO2 observations, global sensitivity analysis of an Euler-Bernoulli beam, and a hierarchical Bayesian model of groundwater pump-test data, are provided to guide users through the process of implementing their own models and leveraging MUQ's UQ algorithms on quasi-realistic applications.

Statement of need

Scientists and engineers are increasingly using physical and statistical models to inform policy, system design, and experiments. Although useful tools, models are inherently error prone and assessing predictive capabilities and robustness requires rigorous uncertainty quantification (UQ). The last decade has seen an explosion in the number and complexity of algorithms for characterizing various sources of uncertainty (e.g., [Kennedy & O'Hagan \(2001\)](#), [Conrad & Marzouk \(2013\)](#), [Giles \(2008\)](#), [Sargsyan et al. \(2019\)](#), [Cotter et al. \(2013\)](#), [Cui et al. \(2016\)](#), [Han & Liu \(2018\)](#), [Detommaso et al. \(2018\)](#)). The complexity of many recent advancements makes it difficult to rigorously compare new algorithms against the current state-of-the-art

^{*}Corresponding Author.

40 and for users to leverage these new tools on practical applications. Likewise, many interesting
41 models are developed, but due to a lack of a common interface they are often not widely
42 used. MUQ aims to reduce the gap between algorithmic research and application by providing
43 a software pipeline between the algorithmic development community and UQ practitioners.
44 The goal is to reduce the costly and error prone practice of reimplementing state-of-the-
45 art techniques, lower the barriers preventing widespread use of cutting-edge techniques, and
46 provide an algorithm-agnostic model interface.

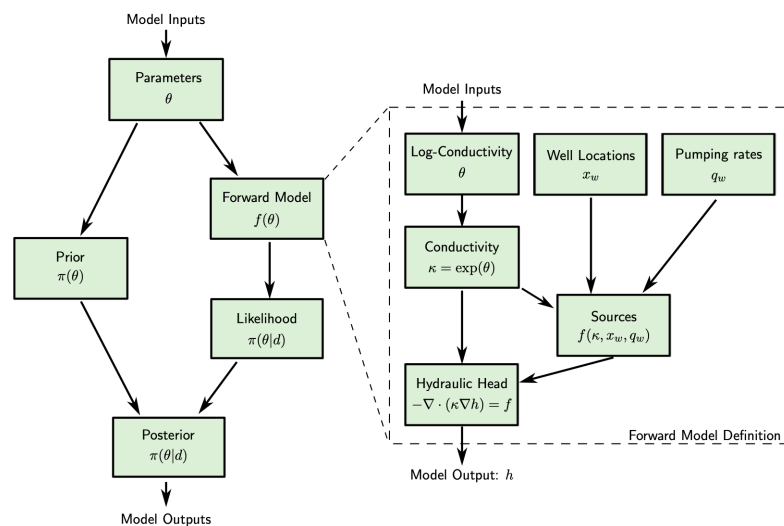


Figure 1: MUQ allows for complicated models to be constructed by connecting model components on a graph. Here is a possible graph for a Bayesian inverse problem built on a model for groundwater flow. MUQ treats each box as a black-box, but if all components can provide derivative information individually, e.g., through adjoint methods, then MUQ can compute gradients, Jacobians, and Hessian actions through the entire graph.

47 While MUQ is capable of solving both forward and inverse UQ problems, its primary focus
48 is on the solution of Bayesian inverse problems with computationally expensive models and
49 potentially high dimensional parameter spaces. This is in contrast to other packages, such as
50 Stan (Carpenter et al., 2017), BUGS (Lunn et al., 2009), or JAGS (Plummer & others, 2003),
51 which are rooted in the statistics community and are not suitable for large-scale models. MUQ
52 also employs a semi-intrusive “gray-box” approach (see Figure Figure 1) that enables efficient
53 gradient calculations, through techniques like the adjoint methods used in PDE constrained
54 optimization, but does not rely on automatic differentiation and does not place any restrictions
55 on how the model is implemented (language, solver, etc...). Other sampling packages, such
56 as PyMC3 (Salvatier et al., 2016), Stan (Carpenter et al., 2017), and tensorflow-probability
57 (Lao et al., 2020), have adopted various probabilistic programming approaches that are more
58 intrusive than MUQ’s gray-box approach and make it more difficult for users to expose effi-
59 cient gradient evaluation techniques for high dimensional problems. MUQ also has a variety of
60 algorithms (e.g., Cotter et al. (2013); Cui et al. (2016)) for tackling discretizations of infinite-
61 dimensional Bayesian inverse problems. MUQ also has several novel MCMC implementations,
62 including multi-level Dodwell et al. (2019) and multi-index MCMC, local approximation algo-
63 rithms Davis et al. (2020), and adaptive transport map MCMC (Parno & Marzouk, 2018).

Acknowledgements

65 We acknowledge additional software contributions from Patrick Conrad and additional support
66 from Youssef Marzouk, Peter Bastian, and Robert Scheichl.

67 This material is based upon work supported by the National Science Foundation under Grant
68 No. ACI-1550487.

69 This material is based upon work supported by the US Department of Energy, Office of
70 Advanced Scientific Computing Research, SciDAC (Scientific Discovery through Advanced
71 Computing) program under awards DE-SC0007099 and DE-SC0021226, for the QUEST and
72 FASTMath SciDAC Institutes.

73 References

- 74 Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., Brubaker,
75 M. A., Guo, J., Li, P., & Riddell, A. (2017). Stan: A probabilistic programming language.
76 *Grantee Submission*, 76(1), 1–32.
- 77 Conrad, P. R., Davis, A. D., Marzouk, Y. M., Pillai, N. S., & Smith, A. (2018). Parallel
78 local approximation MCMC for expensive models. *SIAM/ASA Journal on Uncertainty*
79 *Quantification*, 6(1), 339–373. <https://doi.org/10.1137/16M1084080>
- 80 Conrad, P. R., & Marzouk, Y. M. (2013). Adaptive smolyak pseudospectral approximations.
81 *SIAM Journal on Scientific Computing*, 35(6), A2643–A2670. [https://doi.org/10.1137/](https://doi.org/10.1137/120890715)
82 [120890715](https://doi.org/10.1137/120890715)
- 83 Cotter, S. L., Roberts, G. O., Stuart, A. M., & White, D. (2013). MCMC methods for
84 functions: Modifying old algorithms to make them faster. *Statistical Science*, 424–446.
85 <https://doi.org/10.1214/13-STS421>
- 86 Cui, T., Law, K. J., & Marzouk, Y. M. (2016). Dimension-independent likelihood-informed
87 MCMC. *Journal of Computational Physics*, 304, 109–137. [https://doi.org/10.1016/j.jcp.](https://doi.org/10.1016/j.jcp.2015.10.008)
88 [2015.10.008](https://doi.org/10.1016/j.jcp.2015.10.008)
- 89 Davis, A., Marzouk, Y., Smith, A., & Pillai, N. (2020). Rate-optimal refinement strategies
90 for local approximation MCMC. *arXiv Preprint arXiv:2006.00032*.
- 91 Detommaso, G., Cui, T., Spantini, A., Marzouk, Y., & Scheichl, R. (2018). A stein variational
92 newton method. *arXiv Preprint arXiv:1806.03085*.
- 93 Dodwell, T., Ketelsen, C., Scheichl, R., & Teckentrup, A. (2019). Multilevel markov chain
94 monte carlo. *SIAM Review*, 61, 509–545. <https://doi.org/10.1137/19M126966X>
- 95 Dodwell, T., Ketelsen, C., Scheichl, R., & Teckentrup, A. (2015). *A hierarchical multilevel*
96 *markov chain monte carlo algorithm with applications to uncertainty quantification in*
97 *subsurface flow*. <https://doi.org/10.1137/130915005>
- 98 Giles, M. B. (2008). Multilevel monte carlo path simulation. *Oper. Res.*, 56(3), 607–617.
99 <https://doi.org/10.1287/opre.1070.0496>
- 100 Han, J., & Liu, Q. (2018). Stein variational gradient descent without gradient. *International*
101 *Conference on Machine Learning*, 1900–1908.
- 102 Jakob, W., Rhineland, J., & Moldovan, D. (2017). *pybind11 – seamless operability between*
103 *c++11 and python*.
- 104 Kennedy, M. C., & O'Hagan, A. (2001). Bayesian calibration of computer models. *Journal*
105 *of the Royal Statistical Society: Series B (Statistical Methodology)*, 63(3), 425–464.
- 106 Lao, J., Suter, C., Langmore, I., Chimisov, C., Saxena, A., Sountsov, P., Moore, D., Saurous,
107 R. A., Hoffman, M. D., & Dillon, J. V. (2020). Tfp.mcmc: Modern markov chain monte
108 carlo tools built for modern hardware. *arXiv Preprint arXiv:2002.01184*.
- 109 Lunn, D., Spiegelhalter, D., Thomas, A., & Best, N. (2009). The BUGS project: Evolution,
110 critique and future directions. *Statistics in Medicine*, 28(25), 3049–3067.

- 111 Parno, M. D., & Marzouk, Y. M. (2018). Transport map accelerated markov chain monte
112 carlo. *SIAM/ASA Journal on Uncertainty Quantification*, 6(2), 645–682. [https://doi.](https://doi.org/10.1137/17M1134640)
113 [org/10.1137/17M1134640](https://doi.org/10.1137/17M1134640)
- 114 Plummer, M., & others. (2003). JAGS: A program for analysis of bayesian graphical models
115 using gibbs sampling. *Proceedings of the 3rd International Workshop on Distributed*
116 *Statistical Computing*, 124, 1–10.
- 117 Salvatier, J., Wiecki, T. V., & Fonnesbeck, C. (2016). Probabilistic programming in python
118 using PyMC3. *PeerJ Computer Science*, 2, e55.
- 119 Sargsyan, K., Huan, X., & Najm, H. N. (2019). Embedded model error representation for
120 bayesian model calibration. *International Journal for Uncertainty Quantification*, 9(4).
121 <https://doi.org/10.1615/Int.J.UncertaintyQuantification.2019027384>

DRAFT