



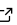
1 lattice-symmetries: A package for working with 2 quantum many-body bases

3 **Tom Westerhout**¹

4 **1** Institute for Molecules and Materials, Radboud University

DOI: [10.21105/joss.03537](https://doi.org/10.21105/joss.03537)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Jed Brown](#) 

Reviewers:

- [@jeffhammond](#)
- [@joselado](#)

Submitted: 26 May 2021

Published: 26 July 2021

License

Authors of papers retain
copyright and release the work
under a Creative Commons
Attribution 4.0 International
License ([CC BY 4.0](#)).

5 Summary

6 Exact diagonalization (ED) is one of the most reliable and established numerical methods of
7 quantum many-body theory. It is precise, unbiased, and general enough to be applicable to a
8 huge variety of problems in condensed matter physics. Mathematically, ED is a linear algebra
9 problem involving a matrix called Hamiltonian. For a system of spin-1/2 particles, the size of
10 this matrix scales exponentially (as $\mathcal{O}(2^N)$) with the number of particles N .

11 Very fast scaling of memory requirements with system size is the main computational challenge
12 of the method. There are a few techniques allowing one to lower the amount of storage used by
13 the Hamiltonian. For example, one can store only the non-zero elements of the Hamiltonian.
14 This is beneficial when the Hamiltonian is sparse which is usually the case in condensed matter
15 physics. One can even take it one step further and avoid storing the matrix altogether by
16 instead computing matrix elements on the fly.

17 A complementary approach to reduce memory requirements is to make use of system symme-
18 tries. For example, many relevant Hamiltonians possess $U(1)$ symmetry which permits one
19 to perform calculations assuming that the number of particles (or number of spins pointing
20 upwards), is fixed. Another example would be translational invariance of the underlying lattice.

21 Although the algorithms for dealing with lattice symmetries are well known ([Sandvik et al.,
22 2010](#)), implementing them remains a non-trivial task. Here we present `lattice-symmetries`,
23 a package providing high-quality and high-performance realization of these algorithms. Instead
24 of writing their own optimized implementation for every system of interest, a domain expert
25 provides system-specific details (such as the number of particles or momentum quantum
26 number) to `lattice-symmetries` and it will automatically construct a reduced Hamiltonian.
27 Dimension of the new Hamiltonian can be multiple orders of magnitude smaller than of the
28 original one.

29 Furthermore, in `lattice-symmetries` the Hamiltonian itself is never stored. Instead, its
30 matrix elements are computed on the fly which reduces the memory requirements even more.
31 Care is taken to keep the implementation generic such that different physical systems can be
32 studied, but without sacrificing performance as we will show in the next section.

33 All in all, `lattice-symmetries` serves as a foundation for building state-of-the-art ED and
34 VMC (Variational Monte Carlo) applications. For example, `SpinED` ([Westerhout, 2020](#)) is an
35 easy-to-use application for exact diagonalization which is built on top of `lattice-symmetri`
36 es and can handle clusters of up to $\mathcal{O}(42)$ spins on a single node.

37 Statement of need

38 Exact diagonalization is an old and well-established method and many packages have been
39 written for it. However, we find that for some reason most state-of-the-art implementations

(Läuchli et al., 2019; Wietek & Läuchli, 2018) are closed-source. There are but three notable open-source projects which natively support spin systems¹: HΦ (Kawamura et al., 2017), SPINPACK (Schulenburg, 2017), and QuSpin (Weinberg & Bukov, 2017).




Figure 1: Performance of matrix-vector products in QuSpin, SPINPACK, and lattice-symmetries. For Heisenberg Hamiltonian on square lattices of different sizes, we measure the time it takes to do a single matrix-vector product. Timings for lattice-symmetries are normalized to 1 to show relative speedup compared to QuSpin, but for reference absolute times in seconds are listed as well. Depending on the system speedups over QuSpin vary between 5 and 22 times, but in all cases lattice-symmetries is significantly faster.

HΦ implements a variety of Hamiltonians, works at both zero and finite temperatures, and supports multi-node computations. However, there are a few points in which lattice-symmetries improves upon HΦ. Firstly, HΦ does not support arbitrary lattice symmetries. Secondly, it uses a custom input file format making it not user-friendly. Finally, since HΦ is an executable, it cannot be used to develop new algorithms.

SPINPACK is another popular solution for diagonalization of spin Hamiltonians. SPINPACK does support user-defined symmetries as opposed to HΦ, but its interface is even less user-friendly. Defining a lattice, Hamiltonian, and symmetries requires writing non-trivial amounts of C code. Finally, SPINPACK is slower than lattice-symmetries as illustrated in Figure 1.

QuSpin is much closer in functionality to lattice-symmetries. It is a high-level Python package which natively supports (but is not limited to) spin systems, can employ user-defined lattice symmetries, and can also perform matrix-free calculations (where matrix elements are computed on the fly). However, QuSpin mostly focuses on ease of use and functionality rather than performance. In lattice-symmetries we follow UNIX philosophy (Salus, 1994) and try to “do one thing but do it well.” Even though lattice-symmetries uses essentially the same algorithms as QuSpin, careful implementation allows us to achieve an order of magnitude speedup as shown in Figure 1.

lattice-symmetries is a library implemented in C++ and C. It provides two interfaces:

- Low-level C interface which can be used to implement ED and VMC applications with focus on performance.

¹There are a few projects targeting fermionic systems and the Hubbard model in particular. Although it is possible to transform a spin Hamiltonian into a fermionic one, it is impractical for large-scale simulations since lattice symmetries are lost in the new Hamiltonian.

63 ▪ A higher-level Python wrapper which allows to easily test and prototype algorithms.

64 We make the library easily installable via Conda package manager.

65 The general workflow is as follows: the user starts by defining a few symmetry generators
 66 (`ls_symmetry/Symmetry` in C/Python) from which `lattice-symmetries` automatically
 67 constructs the symmetry group (`ls_group/Group` in C/Python). The user then proceeds to
 68 constructing the Hilbert space basis (`ls_spin_basis/SpinBasis` in C/Python). For some
 69 applications functionality provided by `SpinBasis` may be sufficient, but typically the user
 70 will construct one (or multiple) quantum mechanical operators (`ls_operator/Operator` in
 71 C/Python) corresponding to the Hamiltonian and various observables. `Operator` can be
 72 efficiently applied to vectors in the Hilbert space (i.e. wavefunctions). Also, in cases when the
 73 Hilbert space dimension is so big that the wavefunction cannot be written down explicitly (as
 74 a list of coefficients), `Operator` can be applied to individual spin configurations to implement
 75 Monte Carlo local estimators.

76 As an example of what can be done with `lattice-symmetries`, we implemented a standalone
 77 application for exact diagonalization studies of spin-1/2 systems: `SpinED`. By combining
 78 `lattice-symmetries` with PRIMME eigensolver (Stathopoulos & McCombs, 2010), it allows
 79 one to treat systems of up to $\mathcal{O}(42)$ sites on a single node. `SpinED` is distributed as a statically-
 80 linked executable — one can download one file and immediately get started with physics. All
 81 in all, it makes large-scale ED more approachable for non-experts.

82 Finally, we would like to note that `lattice-symmetries` and `SpinED` have already been
 83 used in a number of research projects (Astrakhantsev et al., 2021; Bagrov et al., 2020;
 84 Westerhout et al., 2020), and we feel that they could benefit many more. For example,
 85 `lattice-symmetries` is currently even being used to simulate quantum circuits.

86 Acknowledgements

87 The author thanks Nikita Astrakhantsev for useful discussions and testing of (most) new
 88 features. The assistance of Askar Iliasov and Andrey Bagrov with with group-theoretic aspects
 89 of this work is appreciated. This work was supported by European Research Council via
 90 Synergy Grant 854843 – FASTCORR.

91 References

- 92 Astrakhantsev, N., Westerhout, T., Tiwari, A., Choo, K., Chen, A., Fischer, M. H., Carleo,
 93 G., & Neupert, T. (2021). Broken-symmetry ground states of the Heisenberg model on
 94 the pyrochlore lattice. *arXiv Preprint arXiv:2101.08787*.
- 95 Bagrov, A. A., Iliasov, A. A., & Westerhout, T. (2020). Kinetic samplers for neural quantum
 96 states. *arXiv Preprint arXiv:2011.02986*.
- 97 Kawamura, M., Yoshimi, K., Misawa, T., Yamaji, Y., Todo, S., & Kawashima, N. (2017).
 98 Quantum lattice model solver HΦ. *Comput. Phys. Commun.*, 217, 180–192. <https://doi.org/10.1016/j.cpc.2017.04.006>
- 99 Läuchli, A. M., Sudan, J., & Moessner, R. (2019). $S = \frac{1}{2}$ kagome Heisenberg antiferromagnet
 100 revisited. *Phys. Rev. B*, 100(15, 15), 155142. [https://doi.org/10.1103/physrevb.100.](https://doi.org/10.1103/physrevb.100.155142)
 101 155142
- 102 Salus, P. H. (1994). *A quarter century of UNIX*. ACM Press/Addison-Wesley Publishing Co.
- 103 Sandvik, A. W., Avella, A., & Mancini, F. (2010). Computational studies of quantum spin sys-
 104 tems. *AIP Conference Proceedings*, 1297, 135–338. <https://doi.org/10.1063/1.3518900>
- 105

- 106 Schulenburg, J. (2017). SPINPACK. Magdeburg University. [https://www-e.ovgu.de/](https://www-e.ovgu.de/jschulen/spin/)
107 [jschulen/spin/](https://www-e.ovgu.de/jschulen/spin/)
- 108 Stathopoulos, A., & McCombs, J. R. (2010). Primme: Preconditioned iterative multimethod
109 eigensolver—methods and software description. *ACM Trans. Math. Softw.*, 37(2), 1–30.
110 <https://doi.org/10.1145/1731022.1731031>
- 111 Weinberg, P., & Bukov, M. (2017). QuSpin: A python package for dynamics and exact
112 diagonalisation of quantum many body systems part i: Spin chains. *SciPost Phys.*, 2(1,
113 1), 003. <https://doi.org/10.21468/scipostphys.2.1.003>
- 114 Westerhout, T. (2020). SpinED. In *GitHub repository*. [https://github.com/twesterhout/](https://github.com/twesterhout/spin-ed)
115 [spin-ed](https://github.com/twesterhout/spin-ed); GitHub.
- 116 Westerhout, T., Astrakhantsev, N., Tikhonov, K. S., Katsnelson, M. I., & Bagrov, A. A.
117 (2020). Generalization properties of neural network approximations to frustrated magnet
118 ground states. *Nat Commun*, 11(1), 1–8. <https://doi.org/10.1038/s41467-020-15402-w>
- 119 Wietek, A., & Läuchli, A. M. (2018). Sublattice coding algorithm and distributed memory
120 parallelization for large-scale exact diagonalizations of quantum many-body systems. *Phys.*
121 *Rev. E*, 98(3, 3), 033309. <https://doi.org/10.1103/physreve.98.033309>

DRAFT